

CodeAlpha Internship Tasks Summary

Task 1

```
```python
```

## Import libraries

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import os
```

## Load dataset

```
os.chdir(r"C:\Users\Sourav\CodeAlpha\project")
df = pd.read_csv("Iris.csv")
```

## Display basic info

```
print(df.head())
print(df.info())
```

```
...
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 150 entries, 0 to 149
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	150 non-null	int64
1	SepalLengthCm	150 non-null	float64
2	SepalWidthCm	150 non-null	float64
3	PetalLengthCm	150 non-null	float64
4	PetalWidthCm	150 non-null	float64
5	Species	150 non-null	object

```
dtypes: float64(4), int64(1), object(1)
```

```
memory usage: 7.2+ KB
```

```
None
```

```
```python
```

Visualize data

```
sns.pairplot(df, hue="Species") plt.show() ```
```

png

```
```python
```

## Prepare data

```
X = df.drop("Species", axis=1) y = df["Species"]
```

## Split into train and test

```
Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, testsize=0.2,
randomstate=42)
```

## Train model

```
model = RandomForestClassifier() model.fit(Xtrain, ytrain)
```

## Predict and evaluate

```
ypred = model.predict(Xtest) print("Accuracy:", accuracyscore(ytest, ypred))
print(classificationreport(ytest, ypred))
```

```
```
```

Accuracy: 1.0

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Iris-setosa | 1.00 | 1.00 | 1.00 | 10 |
| Iris-versicolor | 1.00 | 1.00 | 1.00 | 9 |
| Iris-virginica | 1.00 | 1.00 | 1.00 | 11 |
| accuracy | | | 1.00 | 30 |
| macro avg | 1.00 | 1.00 | 1.00 | 30 |
| weighted avg | 1.00 | 1.00 | 1.00 | 30 |

```
```python
```

```
```
```

Task 2

```
```python
```

## Import libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import os
```

## Load dataset

```
os.chdir(r"C:\Users\Sourav\CodeAlpha_project")
```

```
df = pd.read_csv("car data.csv")
```

## View the dataset

```
print(df.head())
print(df.info())
```

```
```
```

| | Car_Name | Year | Selling_Price | Present_Price | Driven_kms | Fuel_Type | \ |
|---|----------|------|---------------|---------------|------------|-----------|---|
| 0 | ritz | 2014 | 3.35 | 5.59 | 27000 | Petrol | |
| 1 | sx4 | 2013 | 4.75 | 9.54 | 43000 | Diesel | |
| 2 | ciaz | 2017 | 7.25 | 9.85 | 6900 | Petrol | |
| 3 | wagon r | 2011 | 2.85 | 4.15 | 5200 | Petrol | |
| 4 | swift | 2014 | 4.60 | 6.87 | 42450 | Diesel | |

| | Selling_type | Transmission | Owner |
|---|--------------|--------------|-------|
| 0 | Dealer | Manual | 0 |
| 1 | Dealer | Manual | 0 |
| 2 | Dealer | Manual | 0 |
| 3 | Dealer | Manual | 0 |
| 4 | Dealer | Manual | 0 |

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 301 entries, 0 to 300
```

```
Data columns (total 9 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|---------------|----------------|---------|
| 0 | Car_Name | 301 non-null | object |
| 1 | Year | 301 non-null | int64 |
| 2 | Selling_Price | 301 non-null | float64 |
| 3 | Present_Price | 301 non-null | float64 |
| 4 | Driven_kms | 301 non-null | int64 |
| 5 | Fuel_Type | 301 non-null | object |
| 6 | Selling_type | 301 non-null | object |

```
7    Transmission    301 non-null    object
8    Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
None
```

```
```python
```

## Drop car name

```
df = df.drop("Car_Name", axis=1)
```

## Check categorical features

```
print(df.select_dtypes(include='object').columns)
```

## Encode categorical variables

```
df = pd.getdummies(df, dropfirst=True)
```

## Features and target

```
X = df.drop("SellingPrice", axis=1) y = df["SellingPrice"]
```

## Split into train and test

```
Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, testsize=0.2,
randomstate=42)
```

## Train model

```
model = LinearRegression() model.fit(Xtrain, ytrain)
```

## Predict

```
ypred = model.predict(Xtest)
```

## Evaluation

```
rmse = np.sqrt(meansquarederror(ytest, ypred)) r2 = r2score(ytest, y_pred)
```

```
print("Root Mean Squared Error:", round(rmse, 2)) print("R² Score:",
round(r2, 2))
```

## Actual vs Predicted plot

```
plt.figure(figsize=(8, 6)) sns.scatterplot(x=ytest, y=ypred) plt.xlabel("Actual
Price") plt.ylabel("Predicted Price") plt.title("Actual vs Predicted Car
Prices") plt.grid(True) plt.show()
```

```
```
```

```
Index(['Fuel_Type', 'Selling_type', 'Transmission'], dtype='object')
Root Mean Squared Error: 1.87
R² Score: 0.85
```

```
png
```

```
```python
```

```
```
```

Task 3

```
```python
```

## Import libraries

```
import pandas as pd import seaborn as sns import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression from sklearn.metrics
import meansquarederror, r2score from sklearn.model_selection import
train_test_split import numpy as np import os
```

## Load dataset

```
os.chdir(r"C:\Users\Sourav\CodeAlpha_project")
```

```
df = pd.read_csv('Advertising.csv')
```

## View data

```
print(df.head()) print(df.info()) ```
```

```

 Unnamed: 0 TV Radio Newspaper Sales
0 1 230.1 37.8 69.2 22.1
1 2 44.5 39.3 45.1 10.4
2 3 17.2 45.9 69.3 9.3
3 4 151.5 41.3 58.5 18.5
```

```

4 5 180.8 10.8 58.4 12.9
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
Column Non-Null Count Dtype
--- -
0 Unnamed: 0 200 non-null int64
1 TV 200 non-null float64
2 Radio 200 non-null float64
3 Newspaper 200 non-null float64
4 Sales 200 non-null float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
None

```python

```

Visualize correlations

```

sns.pairplot(df, xvars=['TV', 'Radio', 'Newspaper'], yvars='Sales', kind='reg')
plt.show()

```

Check correlation

```

print(df.corr())

```

Features and Target

```

X = df[['TV', 'Radio', 'Newspaper']] y = df['Sales']

```

Train-Test Split

```

Xtrain, Xtest, ytrain, ytest = train_test_split(X, y, testsize=0.2,
randomstate=42)

```

Train model

```

model = LinearRegression() model.fit(Xtrain, ytrain)

```

Predict

```

ypred = model.predict(Xtest)

```

Evaluate

```
rmse = np.sqrt(meansquarederror(ytest, ypred)) r2 = r2score(ytest, y_pred)

print("Root Mean Squared Error:", round(rmse, 2)) print("R² Score:",
round(r2, 2))
```

Visualize prediction

```
plt.figure(figsize=(8, 6)) sns.scatterplot(x=ytest, y=y_pred) plt.xlabel("Actual
Sales") plt.ylabel("Predicted Sales") plt.title("Actual vs Predicted Sales")
plt.grid(True) plt.show()
```

```

png

|            | Unnamed: 0 | TV       | Radio     | Newspaper | Sales     |
|------------|------------|----------|-----------|-----------|-----------|
| Unnamed: 0 | 1.000000   | 0.017715 | -0.110680 | -0.154944 | -0.051616 |
| TV         | 0.017715   | 1.000000 | 0.054809  | 0.056648  | 0.782224  |
| Radio      | -0.110680  | 0.054809 | 1.000000  | 0.354104  | 0.576223  |
| Newspaper  | -0.154944  | 0.056648 | 0.354104  | 1.000000  | 0.228299  |
| Sales      | -0.051616  | 0.782224 | 0.576223  | 0.228299  | 1.000000  |

Root Mean Squared Error: 1.78  
R² Score: 0.9

png

```python

```